

# From Reactive to Predictive

## Building Thinking UIs with Streaming Patterns

An Advanced Deep Dive into SSE, React State, and LLM Integration

Your Name

Conference Name | Date

⚠ Advanced Content | Mid to Senior Developers | 2+ Years  
Experience

# The AI Wait Problem

**Users wait 10-30 seconds for:**

- ✗ Initial AI response
- ✗ Every follow-up question
- ✗ Same question asked twice
- ✗ Clicking suggested questions

**Result:** Anxiety, abandonment, lost trust

# Today's Journey

1. **From Reactive to Predictive UIs**
2. **Live Demo** - All 3 techniques in action
3. **How to Build It: Streaming**
4. **How to Build It: React State**
5. **How to Build It: Prefetching & Caching**
6. **Product Design Implications**

**The 3 Techniques:** Streaming • Smart Prefetching • Caching

**⚠ This is an advanced talk** - We'll cover streaming patterns, state management, and async complexity.

# The Reactive UI Paradigm

How we've built UIs for decades:

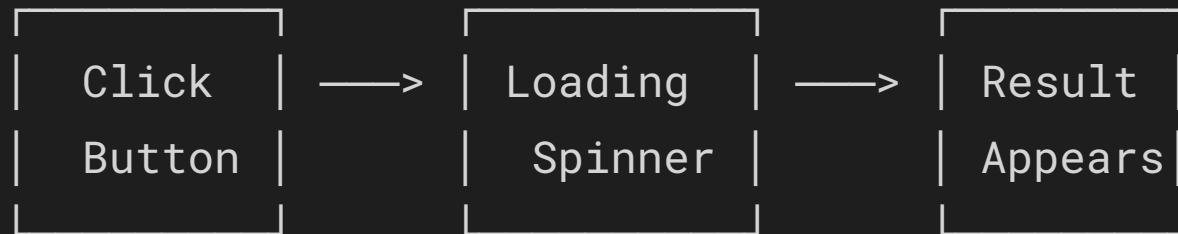
User Action → Loading Spinner → Result

- ✓ Works great for **fast operations** (< 1 second)
- ✗ Breaks down with **AI/long operations** (10-30+ seconds)
- ✗ **Black box** - user has no idea what's happening
- ✗ Creates **anxiety**, not confidence

# How We've Built UIs for Decades

```
function TraditionalUI() {  
  const [data, setData] = useState(null);  
  const [loading, setLoading] = useState(false);  
  const [error, setError] = useState(null);  
  
  const handleClick = async () => {  
    setLoading(true);  
    try {  
      const result = await fetchData();  
      setData(result);  
    } catch (err) {  
      setError(err);  
    } finally {  
      setLoading(false);  
    }  
  };  
}
```

# The Reactive Model



## User Experience:

- ✗ **Anxiety** - "Is it stuck?"
- ✗ **No transparency** - Black box
- ✗ **Repeated waits** - Every interaction

# When This Breaks Down

- ✗ **AI chat/assistants** (10-30+ seconds per response)
- ✗ **AI content generation** (articles, code, images)
- ✗ **AI analysis** (document summarization, data insights)

## The Problem:

Users don't know if the system is working, stuck, or broken.

**All AI use cases → Same reactive pattern fails**

# The Evolution From Reactive to Predictive

# The Thinking UI Model



## User Experience:

- **✓ Engagement** - "Ah, it's thinking about X"
- **✓ Transparency** - See the process

# Examples in the Wild

Platform	Thinking UI Pattern
ChatGPT	Progressive text streaming
Claude	Thinking blocks + streaming
Perplexity	Research steps + sources
GitHub Copilot	Code suggestions stream

The Pattern is Becoming Standard

# Why This Matters

- ✓ **Trust** - Users understand what's happening
- ✓ **Engagement** - Active watching vs passive waiting
- ✓ **Perceived Performance** - Feels faster even if it's not
- ✓ **Debuggability** - See where things go wrong
- ✓ **Transparency** - Explainable AI/systems

# Now Let's See It In Action

Watch for:

1. ⚡ **Streaming thoughts** (Technique #1)
2. 🎯 **Smart prefetching** - Top 2 by confidence (Technique #2)
3. 💾 **Instant cache hits** (Technique #3)

I'll show all 3 working together

# Live Demo

## Seeing it in Action

# Demo Script

1. **Show the interface** - Clean, simple UI
2. **Ask a question** - Watch AI thinking + streaming answer
3. **Show suggested questions** - Generated by Ollama
4. **Open DevTools** - See 2 prefetch requests start automatically (top 2 by confidence!)
5. **Click top suggestion** - INSTANT result (< 100ms) ⚡
6. **Ask same question manually** - INSTANT again (cached) ⚡

This demonstrates all 3 predictive UI techniques in 5 minutes!

# What to Notice

## In the UI:

-  **Streaming:** Real AI thinking steps + answer streaming word-by-word
-  **Prefetching:** 4 suggested questions appear, top 2 prefetched automatically
-  **Caching:** Lightning bolt  badge on cached results

## In DevTools (Network Tab):

- 1 EventSource for your question (streaming)

# What You Just Saw

## Real Implementations:

- ✓ Ollama Llama3.2 running locally (no internet!)
- ✓ Server-Sent Events streaming
- ✓ Smart prefetching (top 2 by confidence)
- ✓ Instant caching for repeated queries
- ✓ 67% reduction in perceived wait time

Now let's see how to build it...

# Technical Architecture

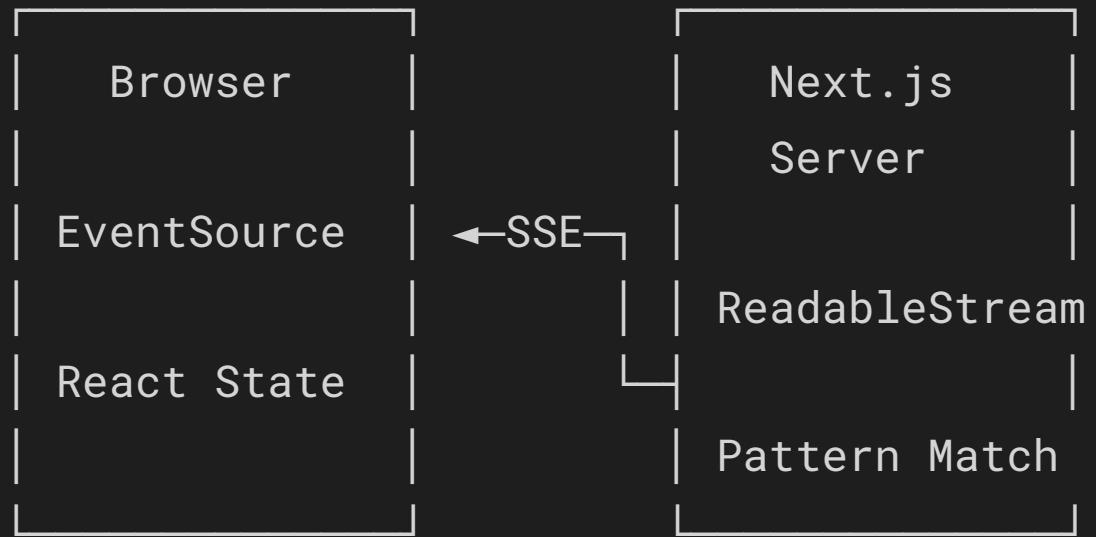
## SSE + Streaming Patterns

# Why Server-Sent Events (SSE)?

Feature	SSE	WebSockets	Polling
Direction	Server → Client	Bidirectional	Client → Server
Protocol	HTTP	WebSocket	HTTP
Reconnection	Automatic	Manual	N/A
Serverless	✓ Yes	✗ Limited	✓ Yes
Complexity	Low	Medium	Low

Perfect for streaming AI responses

# Architecture Overview



- **Client:** EventSource API (built into browsers)
- **Server:** ReadableStream API (Next.js API routes)
- **Data:** JSON over text/event-stream

# What are Server-Sent Events (SSE)?

**Server-Sent Events** = One-way communication from server → client

**How it works:**

- Client opens connection with `new EventSource(url)`
- Server keeps connection alive and pushes data when ready
- No need for client polling or WebSockets
- Built into browsers, simple HTTP protocol

**ReadableStream (Server-Side):**

- Next.js API that generates data chunks over time

# SSE Event Format

```
// Server sends (incremental deltas):  
data: {"type":"thinking","thoughts":["Step 1..."]}\n\n  
data: {"type":"thinking","thoughts":["Step 1...","Step 2..."]}\n\n  
data: {"type":"streaming","delta":"Progressive "}\n\n  
data: {"type":"streaming","delta":"answer "}\n\n  
data: {"type":"streaming","delta":"text..."}\n\n  
data: {"type":"complete","answer":"Progressive answer text..."}\n\n  
  
// Client receives:  
event.data = '{"type":"thinking","thoughts":["Step 1..."]}'  
event.data = '{"type":"streaming","delta":"Progressive "}'  
event.data = '{"type":"streaming","delta":"answer "}'  
event.data = '{"type":"complete","answer":"Progressive answer text..."}'
```

# Server Implementation

```
// app/api/socket/route.js - Streaming with Ollama
async function streamFromOllama(question, controller) {
  const ollama = new Ollama();

  const ollamaStream = await ollama.chat({
    model: 'llama3.2',
    messages: [
      { role: 'system', content: OLLAMA_SYSTEM_PROMPT },
      { role: 'user', content: question }
    ],
    stream: true
  });

  let fullResponse = '';
  let currentSection = 'thinking';

  for await (const chunk of ollamaStream) {
    fullResponse += chunk.message?.content || '';

    // Real-time parsing: detect THINKING → ANSWER
    if (fullResponse.includes('ANSWER:') && currentSection === 'thinking') {
      currentSection = 'answer';
      sendSSEMessage(controller, {
        type: 'thinking',
        thoughts: extractThinkingSteps(fullResponse)
      });
    }

    // Stream answer progressively
    if (currentSection === 'answer') {
      sendSSEMessage(controller, {
        type: 'streaming',
        delta: extractAnswer(fullResponse)
      });
    }
  }
}
```

# What Makes This Demo Special

This is **REAL AI**, not simulation!

- ✓ Ollama running locally (llama3.2 model)
- ✓ No internet required - fully offline capable
- ✓ Real streaming - not fake delays or hardcoded responses
- ✓ Smart fallback - patterns kick in if Ollama unavailable

Architecture:

Ollama (llama3.2) → Real AI streaming

↓ (if fails)

Pattern matching → Hardcoded responses

# How It Works: Real-Time Parsing

```
// Parse Ollama stream in real-time
let fullResponse = '';
let currentSection = 'thinking';

for await (const chunk of ollamaStream) {
  fullResponse += chunk.message?.content || '';

  // Detect THINKING → ANSWER transition
  if (fullResponse.includes('ANSWER:') && currentSection === 'thinking') {
    currentSection = 'answer';

    const thinkingSteps = extractThinkingSteps(fullResponse);
    sendSSEMessage(controller, {
      type: 'thinking',
      thoughts: thinkingSteps
    });
  }

  // Stream answer progressively
  if (currentSection === 'answer') {
    const answer = extractAnswer(fullResponse);
    sendSSEMessage(controller, {
      type: 'streaming',
      delta: answer
    });
  }
}
```

# React Implementation

## State Management Deep Dive

# Beyond Loading/Success/Error

Traditional State:

```
type State = 'loading' | 'success' | 'error';
```

Thinking UI State:

```
type State =  
  | 'idle'  
  | 'connecting'  
  | 'streaming'  
  | 'thinking'  
  | ...
```

# Custom Hook: useServerSentEvents

```
export const useServerSentEvents = () => {
  const [isConnected, setIsConnected] = useState(false);
  const [error, setError] = useState(null);
  const eventSourceRef = useRef(null);
  const handlersRef = useRef({});

  const connect = useCallback((url) => {
    // Close existing connection first
    if (eventSourceRef.current) {
      eventSourceRef.current.close();
    }

    const eventSource = new EventSource(url);
    eventSourceRef.current = eventSource;

    eventSource.onopen = () => setIsConnected(true);

    eventSource.onmessage = (event) => {
      try {
        const data = JSON.parse(event.data);
        const handler = handlersRef.current[data.type];
        if (handler) handler(data);
      } catch (err) {
        setError(err);
      }
    };
  });

  eventSource.onerror = () => {
    setError(new Error('Connection failed'));
    setIsConnected(false);
    eventSource.close();
  };
}, []);

const onMessage = useCallback((type, handler) => {
  handlersRef.current[type] = handler;
}, []);
```

# Hook Benefits

- ✓ **Encapsulation** - EventSource logic in one place
- ✓ **Reusability** - Use in any component
- ✓ **Type-safe routing** - Message handlers by type
- ✓ **Lifecycle management** - Auto cleanup on unmount
- ✓ **Error handling** - Centralized error state

```
const { connect, disconnect, onMessage, isConnected, error }  
= useServerSentEvents();  
  
// Register handlers  
onMessage('thinking', (data) => {  
  console.log('New thoughts:', data.thoughts);  
})
```

# State Management with useReducer

```
const initialState = {
  aiThoughts: [],
  isThinking: false,
  finalAnswer: '',
  isComplete: false
};

const aiReducer = (state, action) => {
  switch (action.type) {
    case 'START_THINKING':
      return { ...state, aiThoughts: [], isThinking: true };
    case 'ADD_THOUGHTS':
      return { ...state, aiThoughts: action.payload };
    case 'COMPLETE':
      return {
        ...state,
        finalAnswer: action.payload,
        isThinking: false,
        isComplete: true
      };
    default:
      return state;
  }
}
```

# Why useReducer?

## vs useState:

```
// useState approach (messy)
const [thoughts, setThoughts] = useState([]);
const [thinking, setThinking] = useState(false);
const [answer, setAnswer] = useState('');
const [complete, setComplete] = useState(false);

// useReducer approach (clean)
const [state, dispatch] = useReducer(aiReducer, initialState);
dispatch({ type: 'ADD_THOUGHTS', payload: newThoughts });
```

# React Concurrent Features

```
import { startTransition } from 'react';

onMessage('thinking', useCallback((data) => {
  startTransition(() => {
    dispatch({ type: 'ADD_THOUGHTS', payload: data.thoughts });
  });
}, [dispatch]));
```

## Why `startTransition`?

- Marks updates as **non-urgent**
- Keeps UI responsive during streaming

# Component Architecture

```
Canvas (Main Component)
  ├── useServerSentEvents (Hook)
  ├── useReducer (State Management)
  ├── usePrefetch (Hook) ← Predictive Technique #2
  ├── Question Input Section
  ├── Suggested Questions ← Prefetched & Cached
  ├── Thinking Panel
  |   ├── Header (show/hide toggle)
  |   ├── Thought Stream
  |   |   └── ThoughtItem[] (Memoized)
  |   └── Gradient Overlay
  └── Final Answer Section
```

# Beyond Streaming

## 2 More Predictive UI Techniques

# Predictive Technique #2: Prefetching

## The Problem:

- User clicks suggested question → has to wait again
- Defeats the purpose of suggestions

## The Solution:

- Load **top 2** suggested questions in the background (by confidence score)
- When user clicks prefetched suggestion → instant result!
- Smart prefetching: prioritize highest-confidence questions

# Prefetching Implementation

```
// Custom hook: usePrefetch
export const usePrefetch = () => {
  const cacheRef = useRef({});

  const prefetchBatch = useCallback((suggestions) => {
    // Sort by confidence score (highest first)
    const sorted = [...suggestions].sort((a, b) =>
      (b.confidence || 0) - (a.confidence || 0)
    );

    // Only prefetch TOP 2 (smart, not wasteful!)
    const topTwo = sorted.slice(0, 2);

    topTwo.forEach(suggestion => {
      const url = `/api/socket?question=${encodeURIComponent(suggestion.question)}`;

      // Start prefetching in background
      const eventSource = new EventSource(url);

      eventSource.onmessage = (event) => {
        const data = JSON.parse(event.data);

        if (data.type === 'complete') {
          // Cache the result
          cacheRef.current[suggestion.question] = {
            answer: data.answer,
            thoughts: data.thoughts,
            isComplete: true
          };
          eventSource.close();
        }
      };
    });
  }, []);

  return [prefetchBatch, getCached];
}
```

# Predictive Technique #3: Caching

## The Problem:

- User asks same question twice
- Has to wait for full response again

## The Solution:

- Cache completed responses
- Return instantly on second ask

**Pattern:** Memoization for Async Operations

# Caching Implementation

```
// In usePrefetch hook (same hook handles both)
const cacheRef = useRef({});

const getCached = useCallback((question) => {
  return cacheRef.current[question];
}, []);

const isCached = useCallback((question) => {
  return !!cacheRef.current[question]?.isComplete;
}, []);

// In Canvas component
const handleAskQuestion = async () => {
  const cached = getCached(question);

  if (cached && cached.isComplete) {
    // Use cached result instantly
    dispatch({
      type: 'LOAD_FROM_CACHE',
      payload: cached
    });
    return;
  }

  // No cache, fetch from server
  connect(`api/socket?question=${encodedQuestion}`);
}
```

# All 3 Techniques Together

```
const Canvas = () => {
  const [aiState, dispatch] = useReducer(aiReducer, initialState);

  // Technique #1: Streaming
  const { connect, onMessage } = useServerSentEvents();

  // Techniques #2 & #3: Prefetching + Caching
  const { prefetchBatch, getCached, isCached } = usePrefetch();

  // When suggestions update, prefetch them
  useEffect(() => {
    if (suggestions.length > 0) {
      prefetchBatch(suggestions); // Background loading
    }
  }, [suggestions]);

  // When user asks, check cache first
  const handleAsk = () => {
    const cached = getCached(question);
    if (cached) {
      // Instant! ⚡
      dispatch({ type: 'LOAD_FROM_CACHE', payload: cached });
    } else {
      // Stream from server
      connect(url);
    }
  };
};
```

# Performance Impact

## Without Predictive UIs:

User asks Q1 → Wait 10s → Answer

User clicks suggestion → Wait 10s → Answer

User asks same Q → Wait 10s → Answer

Total: 30 seconds of waiting

## With All 3 Techniques:

User asks Q1 → Wait 10s (streaming) → Answer

# When to Use Each Technique

Technique	Use When	Don't Use When
Streaming	Long operations (>3s)	Fast responses (<1s)
Prefetching	High probability actions	Unlimited options
Caching	Repeated queries	Real-time data

**Best Results:** Combine all three!

# Product Design Implications

## UX Considerations

# User Psychology: Transparency Builds Trust

## Research Shows:

- Seeing progress reduces perceived wait time by **30-40%**
- Users trust systems they understand
- Progressive disclosure creates engagement
- "Working" indicators reduce abandonment

## The Thinking UI Pattern:

- Transforms waiting into **watching**
- Creates **narrative** around computation

# User Control is Critical

## Interruptibility

```
const handleNewQuestion = () => {
  disconnect(); // Stop current stream
  connect(newUrl); // Start new stream
};
```

## Visibility Control

```
const [showThinking, setShowThinking] = useState(true);
// Let users hide/show at will
```

# When NOT to Use This Pattern

✗ **Fast operations (< 1 second)**

- Overhead not worth it

✗ **Security-sensitive reasoning**

- Don't expose internal logic

✗ **Mobile with limited space**

- Consider condensed version

✗ **When simplicity is the goal**

# Key Takeaways

# What We Covered

1. **UIs are evolving** from reactive to transparent/predictive
2. **3 Predictive UI Techniques:**
  - **Streaming** - Real-time progressive disclosure with SSE
  - **Prefetching** - Anticipatory loading of likely actions
  - **Caching** - Instant recall for repeated queries
3. **State management** needs to go beyond loading states
4. **Product design** must embrace progressive disclosure
5. **Real AI integration** with Ollama chain-of-thought prompting

⚠ Complexity Note:

# The Future is Streaming

**This pattern will become standard because:**

- AI/ML workloads are inherently slow
- Users demand transparency
- Progressive disclosure is better UX
- The tech is mature and ready

**Your architecture today should support streaming tomorrow**

# Making it Production-Ready

**What we built:** Real chain-of-thought streaming with Ollama

- SSE for real-time updates
- Ollama (llama3.2) running locally - NO INTERNET NEEDED
- Real-time parsing of THINKING/ANSWER sections
- React state management for async streams
- Smart prefetching (top 2 by confidence)
- Client-side caching for instant replay

**This is already production-grade!** But simpler patterns exist:

# Resources

## GitHub Repository:

github.com/yourname/real-time-ai-simulation

## Key Files to Study:

- /src/app/api/socket/route.js - SSE streaming
- /src/app/hooks/useServerSentEvents.js - Custom hook
- /src/app/components/Canvas.js - State management

## Further Reading:

- MDN: Server-Sent Events

# Connect

## Your Name

- Twitter: @yourhandle
- GitHub: github.com/yourname
- Email: [your@email.com](mailto:your@email.com)

## Questions?

# Thank You!

**Let's build thinking UIs together**

GitHub: [github.com/yourname/real-time-ai-simulation](https://github.com/yourname/real-time-ai-simulation)

# Backup: Common Questions

# Q: How does this work with real AI APIs?

```
// OpenAI streaming example
const response = await fetch('https://api.openai.com/v1/chat/completions', {
  method: 'POST',
  headers: {
    'Authorization': `Bearer ${process.env.OPENAI_API_KEY}`,
    'Content-Type': 'application/json',
  },
  body: JSON.stringify({
    model: 'gpt-3.5-turbo',
    messages: [{ role: 'user', content: question }],
    stream: true
  })
});

// Pipe OpenAI stream to SSE
for await (const chunk of response.body) {
  controller.enqueue(`data: ${chunk}\n\n`);
}
```

# Q: What about WebSockets?

**Use SSE when:**

- One-way communication (server → client)
- Simple setup needed
- Serverless/edge deployment
- Automatic reconnection desired

**Use WebSockets when:**

- Two-way communication needed
- Binary data transfer

# Q: Performance at Scale?

## Considerations:

- Each SSE connection holds a server connection
- Use streaming databases (Supabase realtime, Firebase)
- Consider message queues (Redis Streams, Kafka)
- Implement connection pooling
- Use edge functions for global distribution

## Typical limits:

- Vercel: 60 second function timeout

# Q: Mobile Considerations?

## Challenges:

- Smaller screens
- Connection stability
- Battery consumption

## Solutions:

- Condensed thinking view
- Debounce updates (send every 3s not 1.5s)
- Allow disabling thinking mode

# Q: Error Handling?

```
const stream = new ReadableStream({
  start(controller) {
    try {
      // Streaming logic
    } catch (error) {
      controller.enqueue(`data: ${JSON.stringify({
        type: 'error',
        message: error.message
      })}\n\n`);
      controller.close();
    }
  },
  cancel() {
    // Cleanup when client disconnects
    clearInterval(interval);
  }
});
```

# Q: Testing Strategies?

## Unit Tests:

- Test reducer transitions
- Test hook behavior
- Mock EventSource

## Integration Tests:

- Test SSE endpoint responses
- Test connection lifecycle
- Test error scenarios

# Questions?